

Introduction To MicroC/OS-II

Chi-Yang Hung

confidential

Home Automation, Networking and Entertainment Lab.
Dept. of Computer Science and Information Engineering
hanel.csie.ncku.edu.tw

HANEL

OUTLINES

- ❑ Introduction
- ❑ Real-Time Systems Concepts
- ❑ Kernel Structure
- ❑ Porting MicroC/OS-II

Introduction

Home Automation, Networking and Entertainment Lab.

CSIE @ National Cheng Kung University.



Introduction to RTOS

- ❑ What is a Real-Time System?
- ❑ What is a Embedded System?
- ❑ What is a Real-Time Operating System?

What is a Real-Time System?

- ❑ A system enforcing timing constraints, e.g. Avionics, Missile Control, ...
 - The correctness of the system depends not only on the logical result of the computation, but also on the **time** at which the results are produced.
 - Real-Time vs. High Performance
- ❑ What is a timing constraint?
 - A constraint of timing requirements, e.g. deadline, ready time, response time ...

Real-Time System Examples

- ❑ Industrial and automation system
- ❑ Computer networking system
- ❑ Gaming and multimedia
- ❑ Medical instrument and devices
- ❑ Financial transaction applications
- ❑ Military defense system
- ❑ Security monitoring and response system
- ❑ Data acquisition system
- ❑ Machine vision/translation system
- ❑ ...

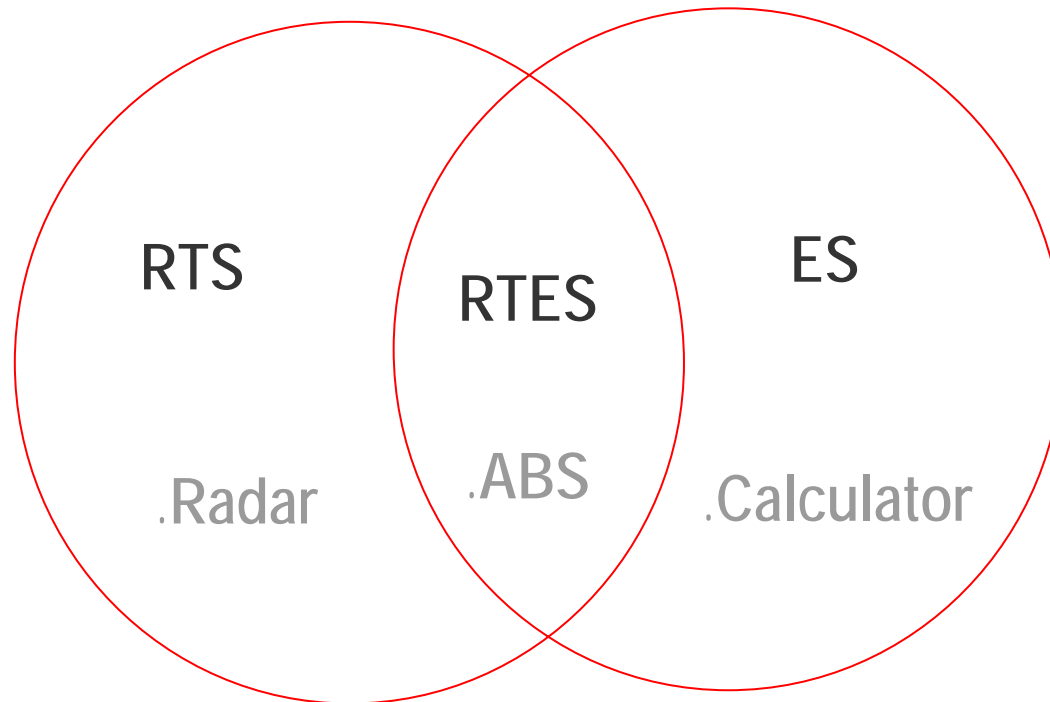
What is a Embedded System?

- ❑ A system designed to perform a **specific** function, e.g. eBook, PDA, eWatch, ...
 - A combination of computer hardware and software, and perhaps additional mechanical or other parts.
 - Embedded vs. General-Purpose
- ❑ What is a specific function?
 - Is there a limitation?
- ❑ Robust, Low-power, Small, ...

Embedded System Examples

- ❑ Computer peripherals
 - Keyboard, Mouse, ...
- ❑ Information Appliances
 - Set-Up Boxes, WebTV, ...
- ❑ Monitors and Sensors
 - Fire Alarm, Heartbeat Detector, ...
- ❑ Controllers in Electronics
 - Refrigerator, Air Conditioner, ...
- ❑ Communication Devices
 - Hub, Router, ...
- ❑ ...

Real-Time Systems vs. Embedded Systems



What is a Real-Time Operating System?

- ❑ What is a real-time operating system?
 - An operating system enforcing timing constraints, Lynx, pSOS, VxWorks, eCOS, uCLinux, LynxOS, RTLinux, KURT, **uC/OS-II**, ...

Real-Time Systems Concepts

Home Automation, Networking and Entertainment Lab.

CSIE @ National Cheng Kung University.



Real-Time Systems Concepts

- ❑ Multitasking
- ❑ Kernel
- ❑ Scheduling
- ❑ Mutual Exclusion
- ❑ Synchronization
- ❑ Interrupt

Multitasking

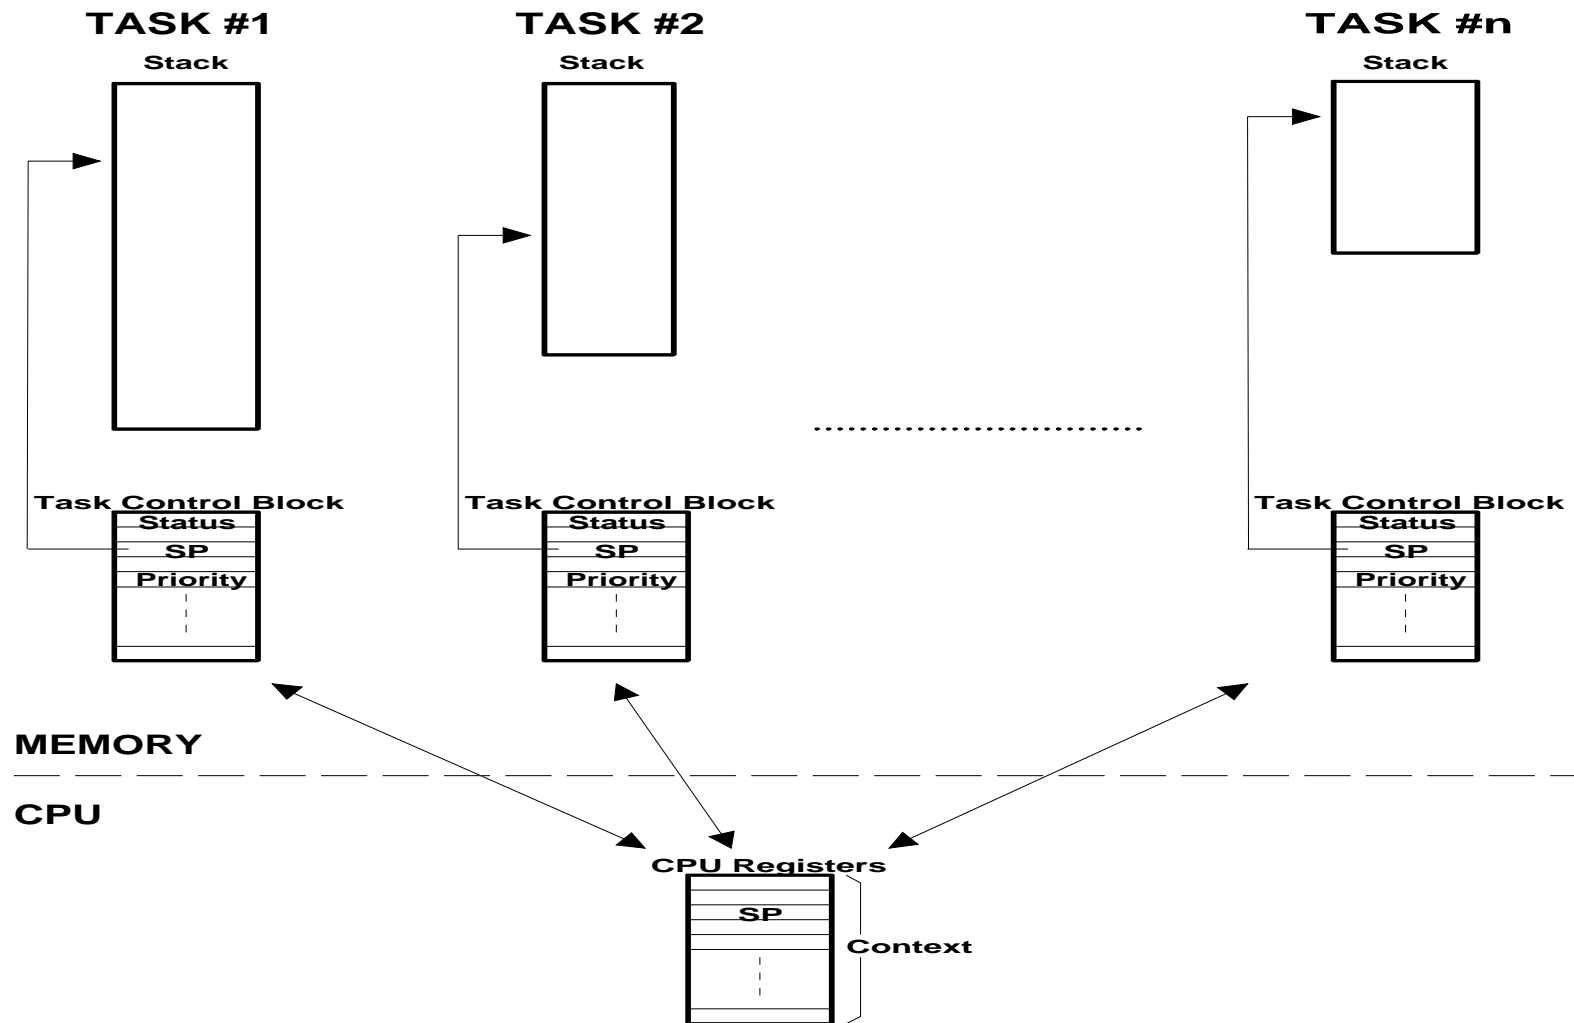
- ❑ Multitasking

- A process of scheduling and switching CPU between several tasks.

- ❑ Related issues

- Context Switch (Task Switch)
- Resource Sharing
- Critical Section

Multitasking(Cont.)



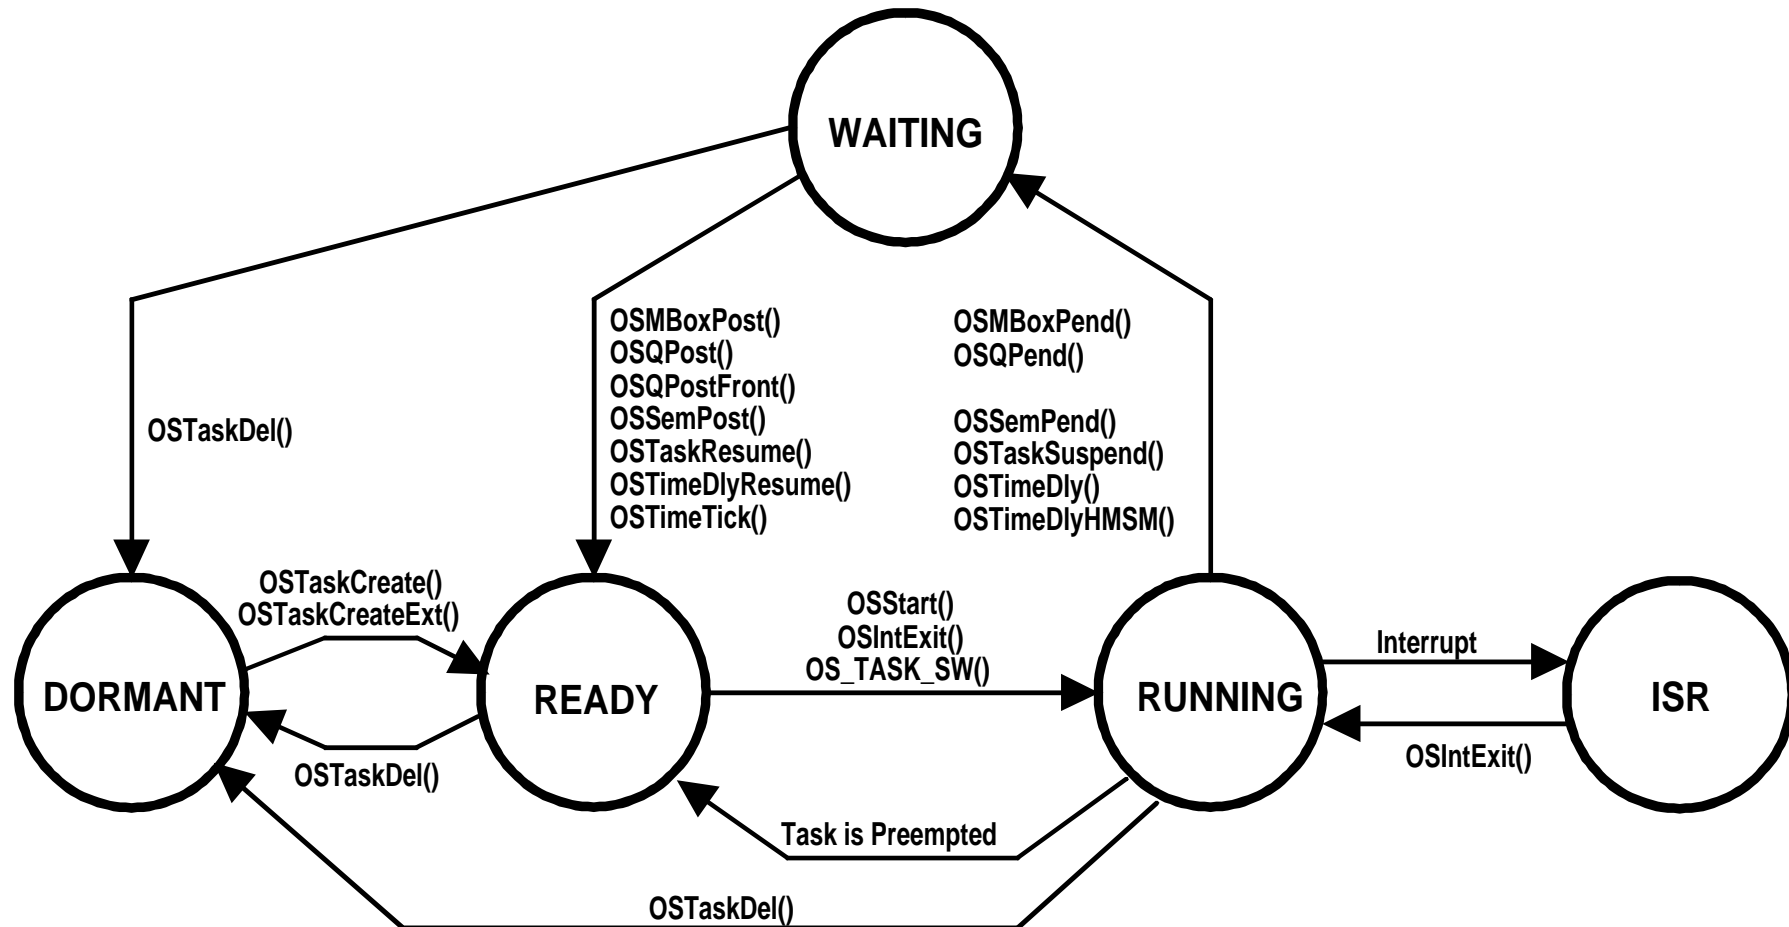
Tasks

- ❑ A task, or called a thread, a simple program that thinks it has the CPU all to itself.
- ❑ Each task is assigned a **priority**, its own set of CPU registers, and its own stack area .
- ❑ Each task typically is an infinite loop that can be in any one of five states.
 - Ready, Running, Waiting, ISR, Dormant

Tasks(Cont.)

- ❑ The DORMANT state corresponds to a task which resides in memory but has not been made available to the multitasking kernel.
- ❑ A task is READY when it can execute but its priority is less than the currently running task.
- ❑ A task is RUNNING when it has control of the CPU.
- ❑ A task is WAITING when it requires the occurrence of an event (waiting for an I/O operation to complete, a shared resource to be available, a timing pulse to occur, time to expire etc.).

Task States



Kernel

- ❑ Kernel is the part of a multitasking system responsible for the management of tasks.
- ❑ Context switching is the fundamental service of a kernel.

Non-Preemptive Kernel

v.s

Preemptive Kernel

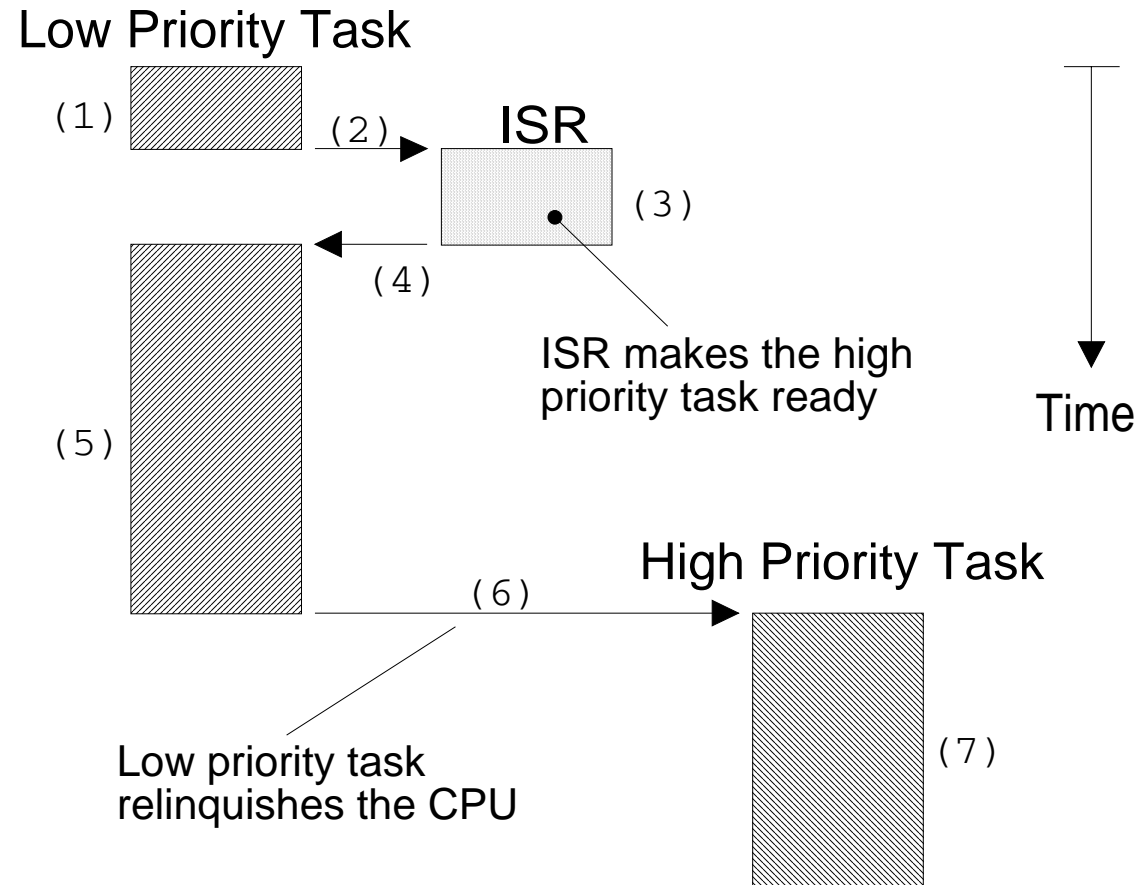
Non-Preemptive Kernel

- ❑ A non-preemptive kernel allows each task to run until it voluntarily gives up control of the CPU.
- ❑ An ISR can make a higher priority task **ready** to run, but the ISR always returns to the interrupted task.
- ❑ One of the advantages of a non-preemptive kernel is that interrupt latency is typically low.
- ❑ Linux 2.4 is non-preemptive.
- ❑ Linux 2.5 is to be preemptive.

Non-Preemptive Kernel(Cont.)

- ❑ At the task level, non-preemptive kernels can also use non-reentrant functions.
- ❑ Another advantage of non-preemptive kernels is the lesser need to guard shared data through the use of semaphores.
- ❑ The most important drawback of a non-preemptive kernel is responsiveness.
- ❑ Very few commercial kernels are non-preemptive.

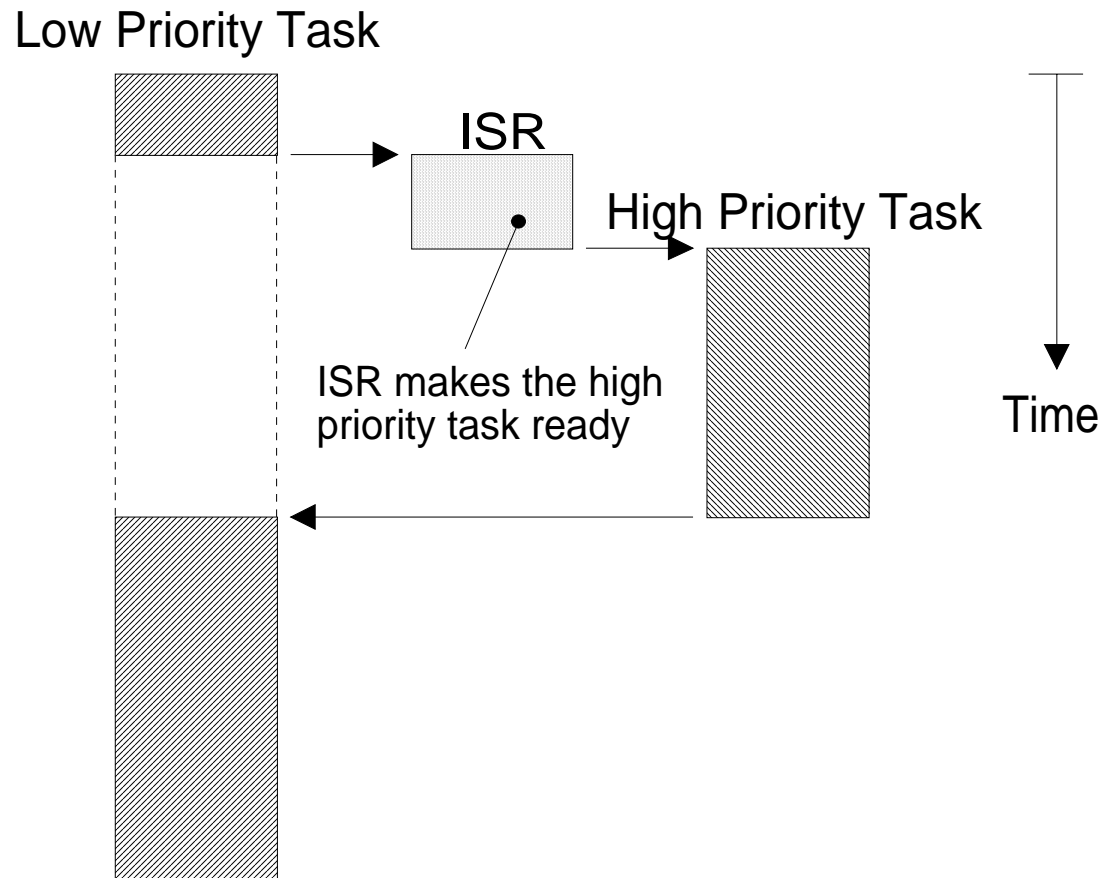
Non-Preemptive Kernel(Cont.)



Preemptive Kernel

- ❑ A preemptive kernel is used when system responsiveness is important.
- ❑ A preemptive kernel always executes the highest priority task that is ready to run.
- ❑ μ C/OS-II and most commercial real-time kernels are **preemptive**.
- ❑ Much better response time.
- ❑ Should not use non-reentrant functions, unless the functions are mutual exclusive.

Preemptive Kernel(Cont.)



Function Reentrancy

- ❑ A reentrant function is a function that can be used by more than one task without fear of data corruption.
- ❑ Reentrant functions either use local variables or protected global variables.
 - `OS_ENTER_CRITICAL()`
 - `OS_EXIT_CRITICAL()`

Function Reentrancy(Cont.)

Non-Reentrant Function

```
static int Temp;
```

```
void swap(int *x, int *y)
{
    Temp = *x;
    *x = *y;
    *y = Temp;
}
```

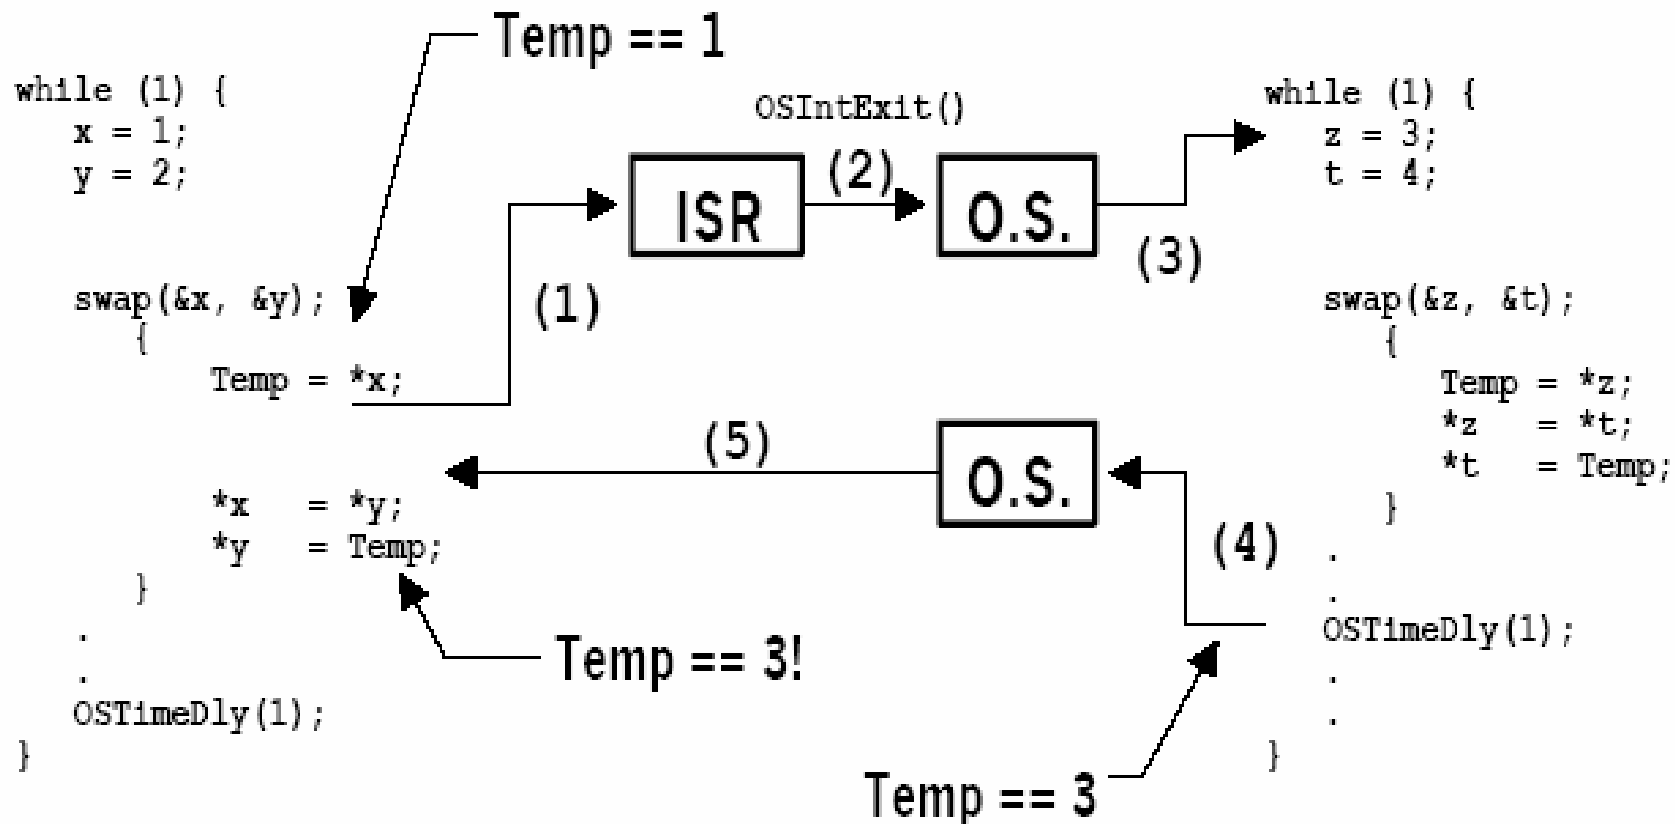
Reentrant Function

```
void strcpy(char *dest, char *src)
{
    while (*dest++ = *src++) {
        ;
    }
    *dest = NULL;
}
```

Function Reentrancy(Cont.)

LOW PRIORITY TASK


HIGH PRIORITY TASK




Scheduling

- ❑ Task Priority Assignment
 - Static priority
 - Rate Monotonic (RM)
 - Dynamic priority
 - Earliest-Deadline First (EDF)

Priority Inversion Problem



Priority Inheritance



Mutual Exclusion

- ❑ Protected shared data of processes.
- ❑ Exclusive access implementation
 - Disabling and enabling interrupts
 - Test-and-Set
 - Disabling and enabling scheduler
 - Semaphores
 - Binary Semaphore
 - Counting Semaphore
- ❑ Deadlock – set timeout for a semaphore

Semaphore

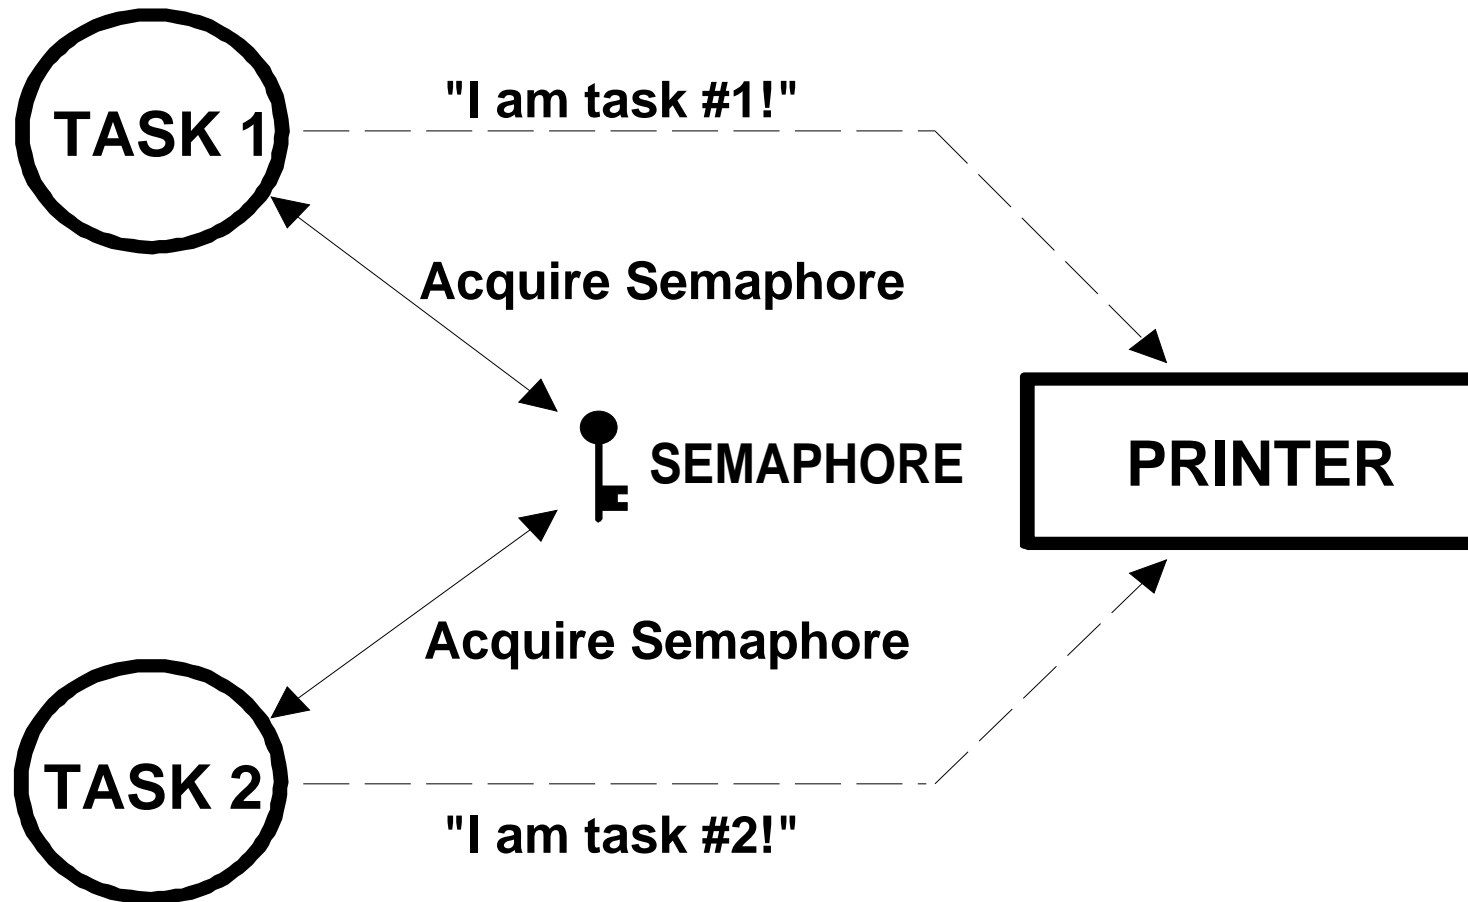
- There are generally only three operations that can be performed on a semaphore: INITIALIZE (also called *CREATE*), WAIT (also called *PEND*), and SIGNAL (also called *POST*).

```
OS_EVENT *SharedDataSem;

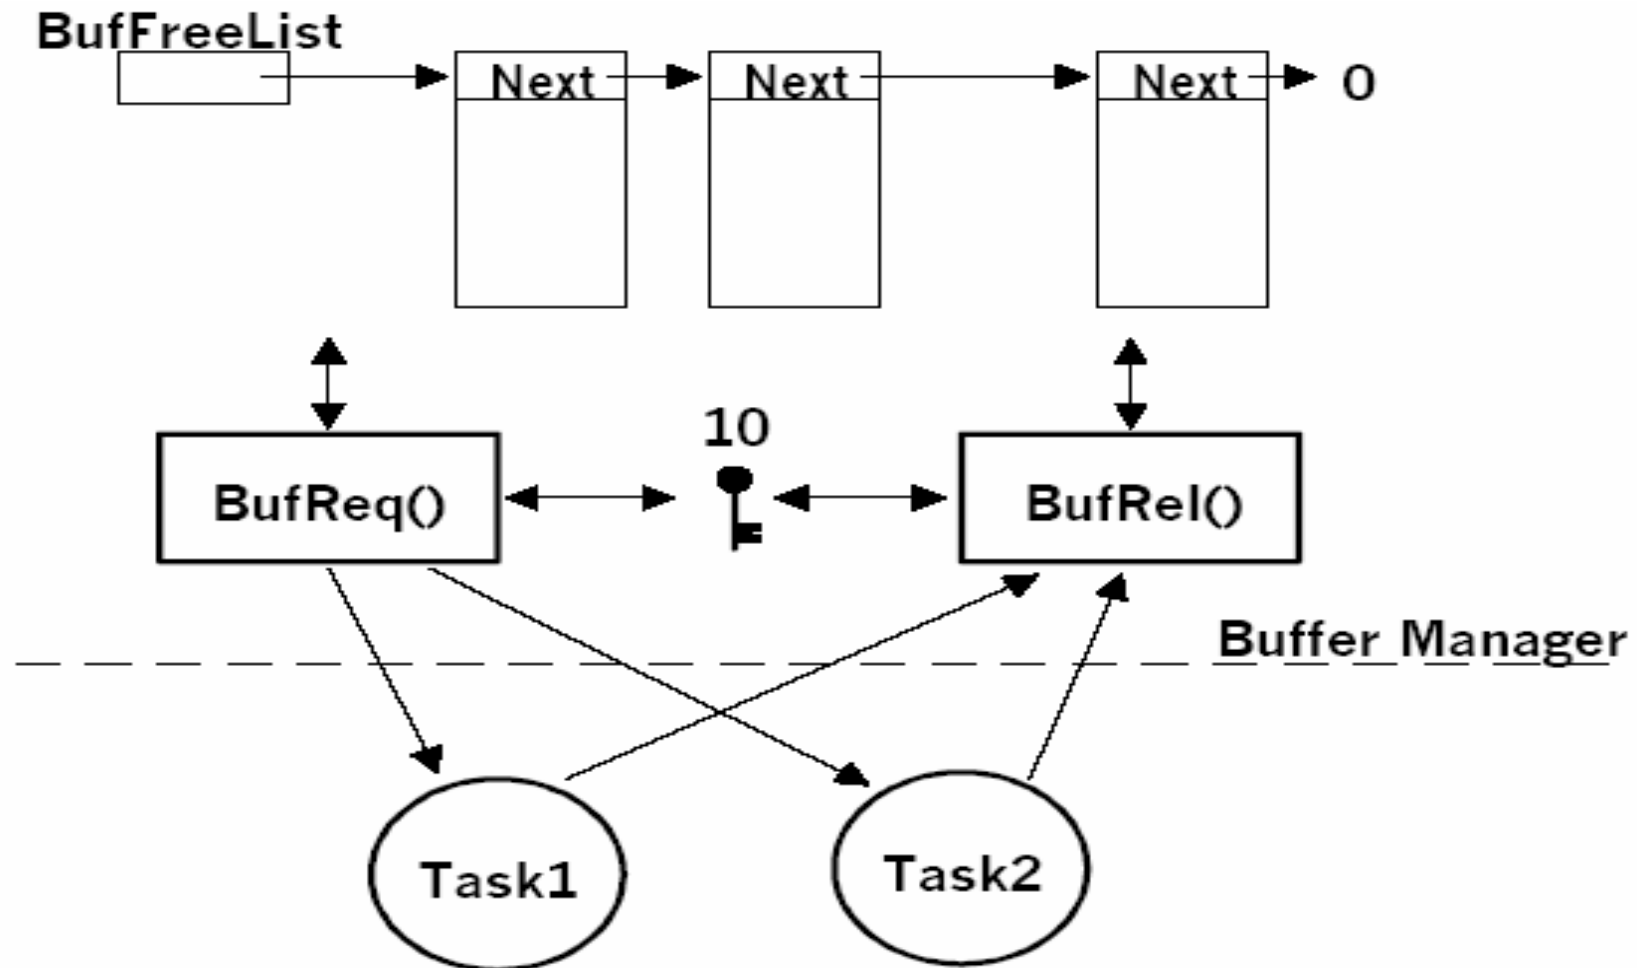
void Function (void)
{
    INT8U err;

    OSSemPend(SharedDataSem, 0, &err);
    .
    .    /* You can access shared data in here (interrupts are recognized) */
    .
    OSSemPost(SharedDataSem);
}
```

Using Binary Semaphore




Using Counting Semaphore




Synchronization

- ❑ Synchronization mechanism is used between tasks or task to ISR.
- ❑ Unilateral rendezvous
 - A task can be synchronized with an ISR, or another task when no data is being exchanged, by using a semaphore.
- ❑ Bilateral rendezvous
 - Two tasks can synchronize their activities by using two semaphores.

Unilateral rendezvous



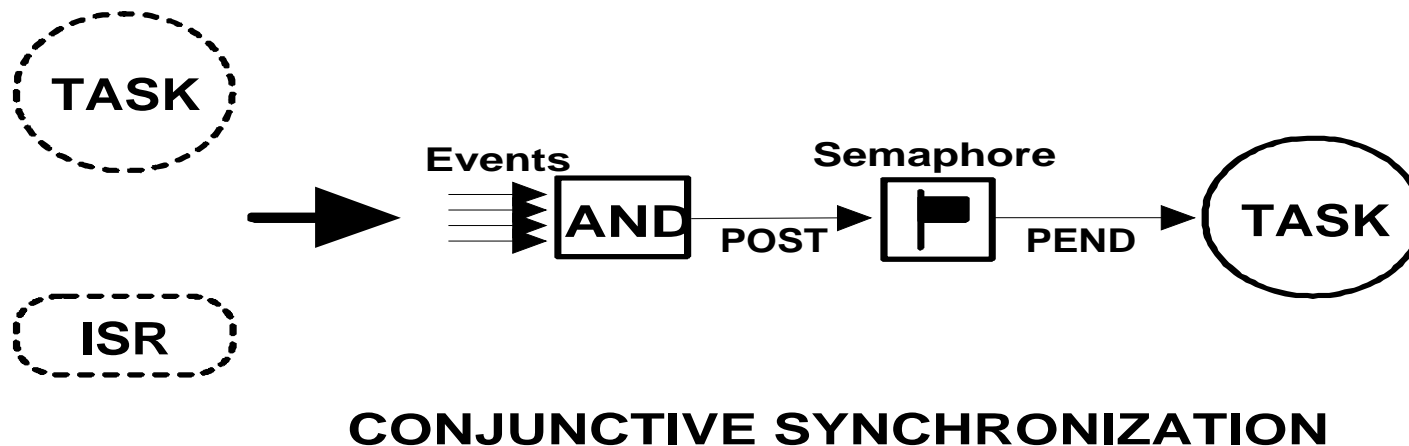

Bilateral rendezvous



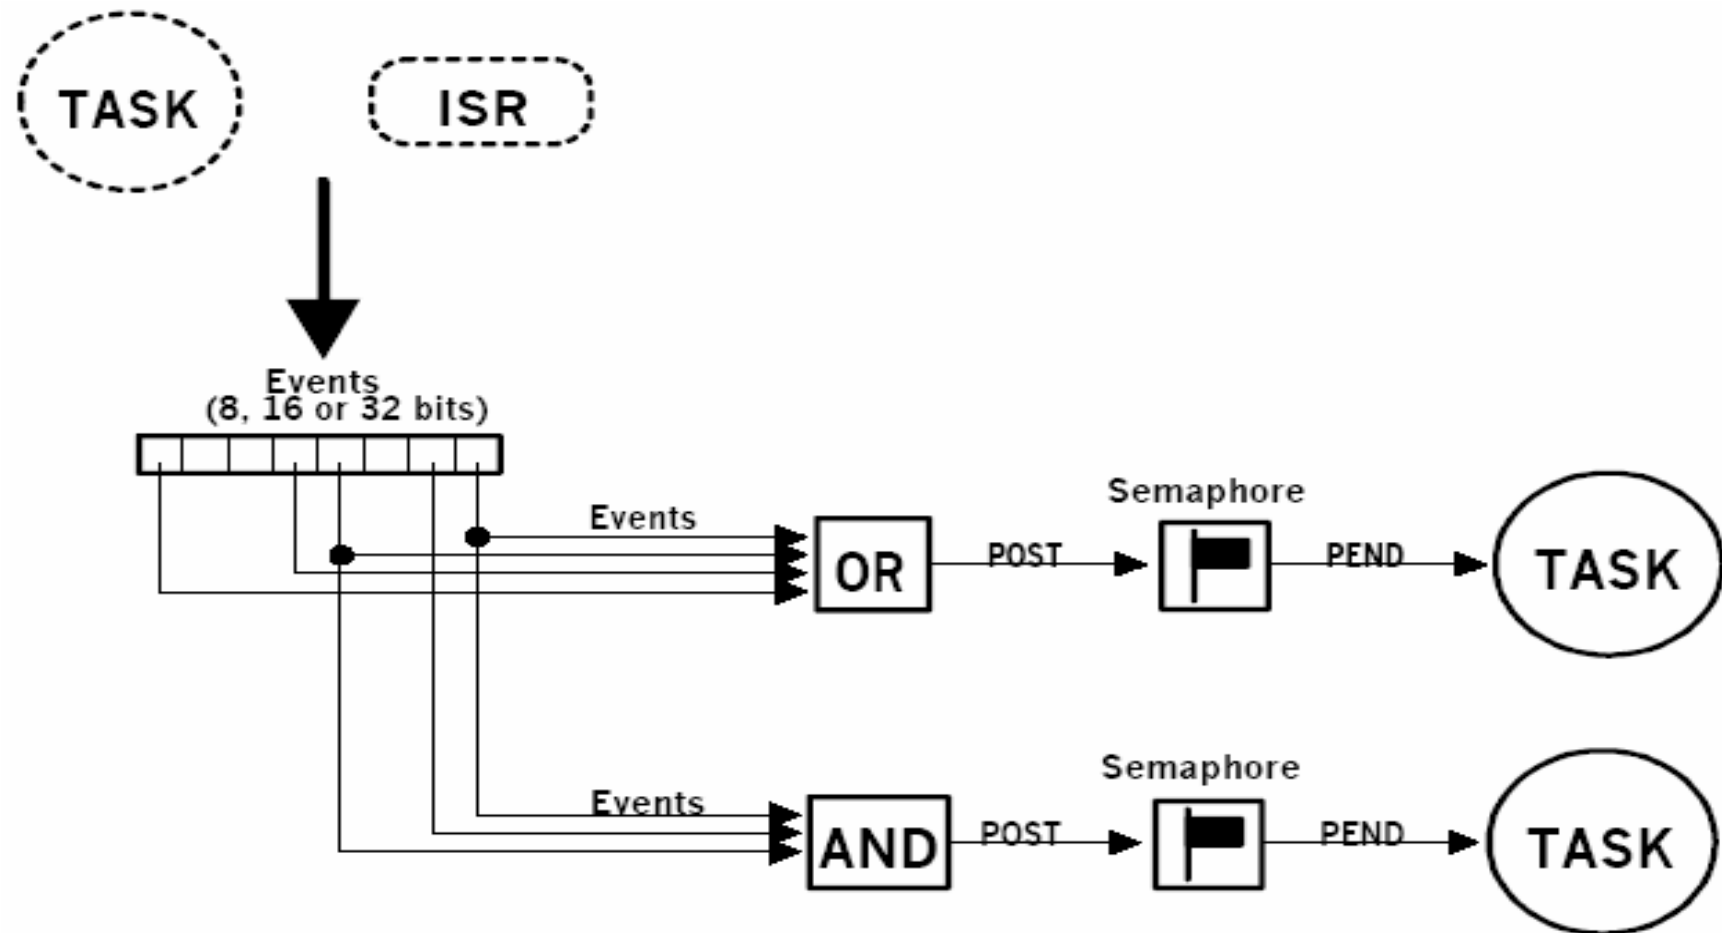
Event Flags

- ❑ The task can be synchronized when **any** of the events have occurred. This is called *disjunctive synchronization* (logical OR).
- ❑ A task can also be synchronized when **all** events have occurred. This is called *conjunctive synchronization* (logical AND).

Event Flags(Cont.)



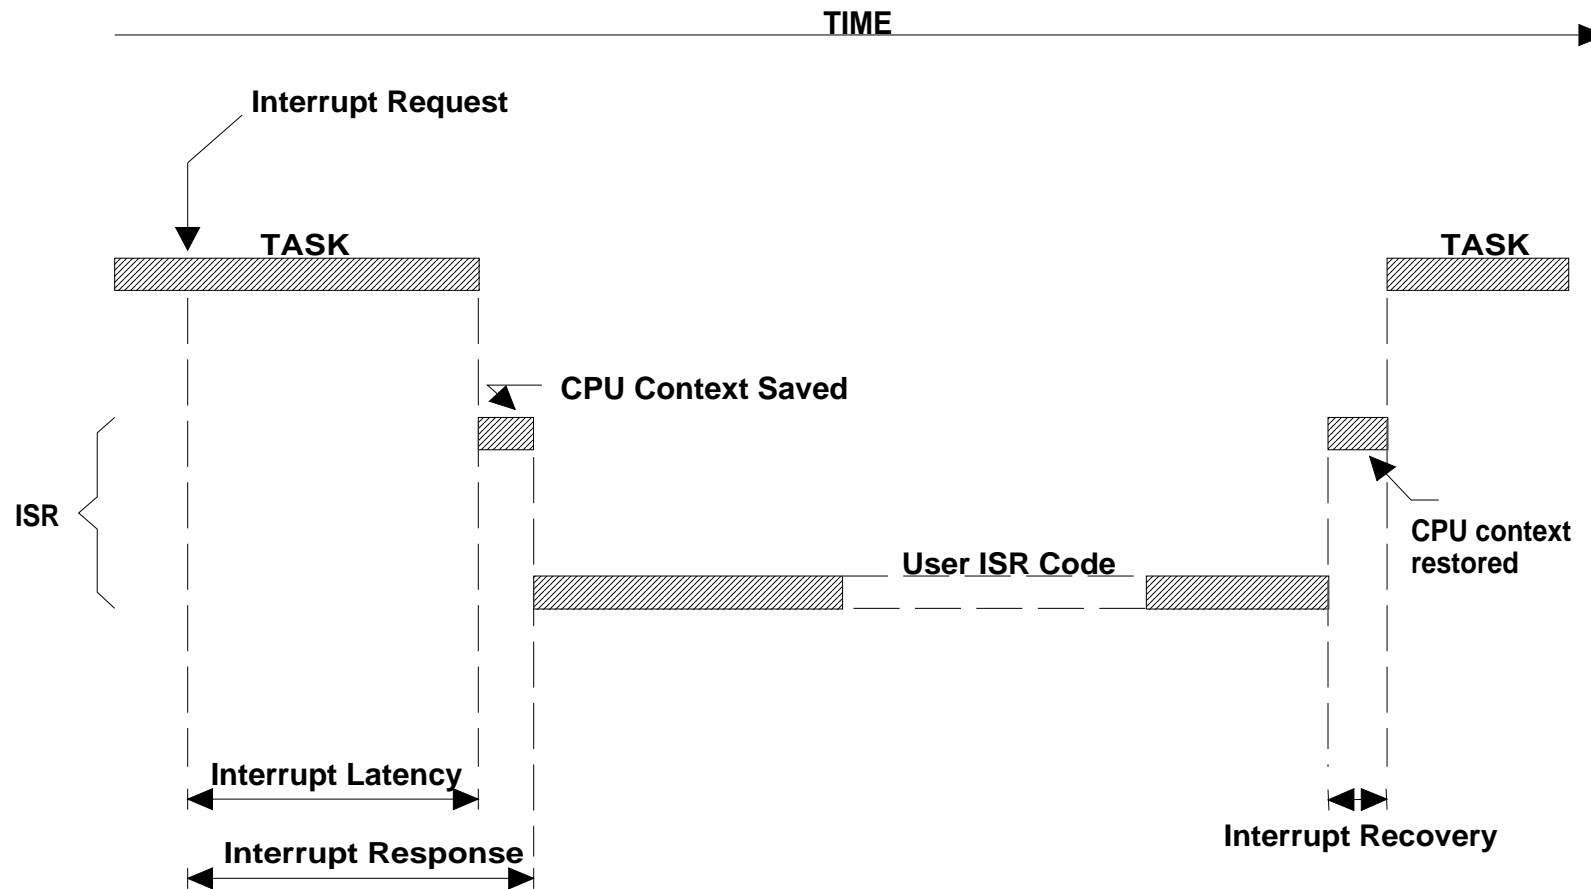
Event Flags(Cont.)



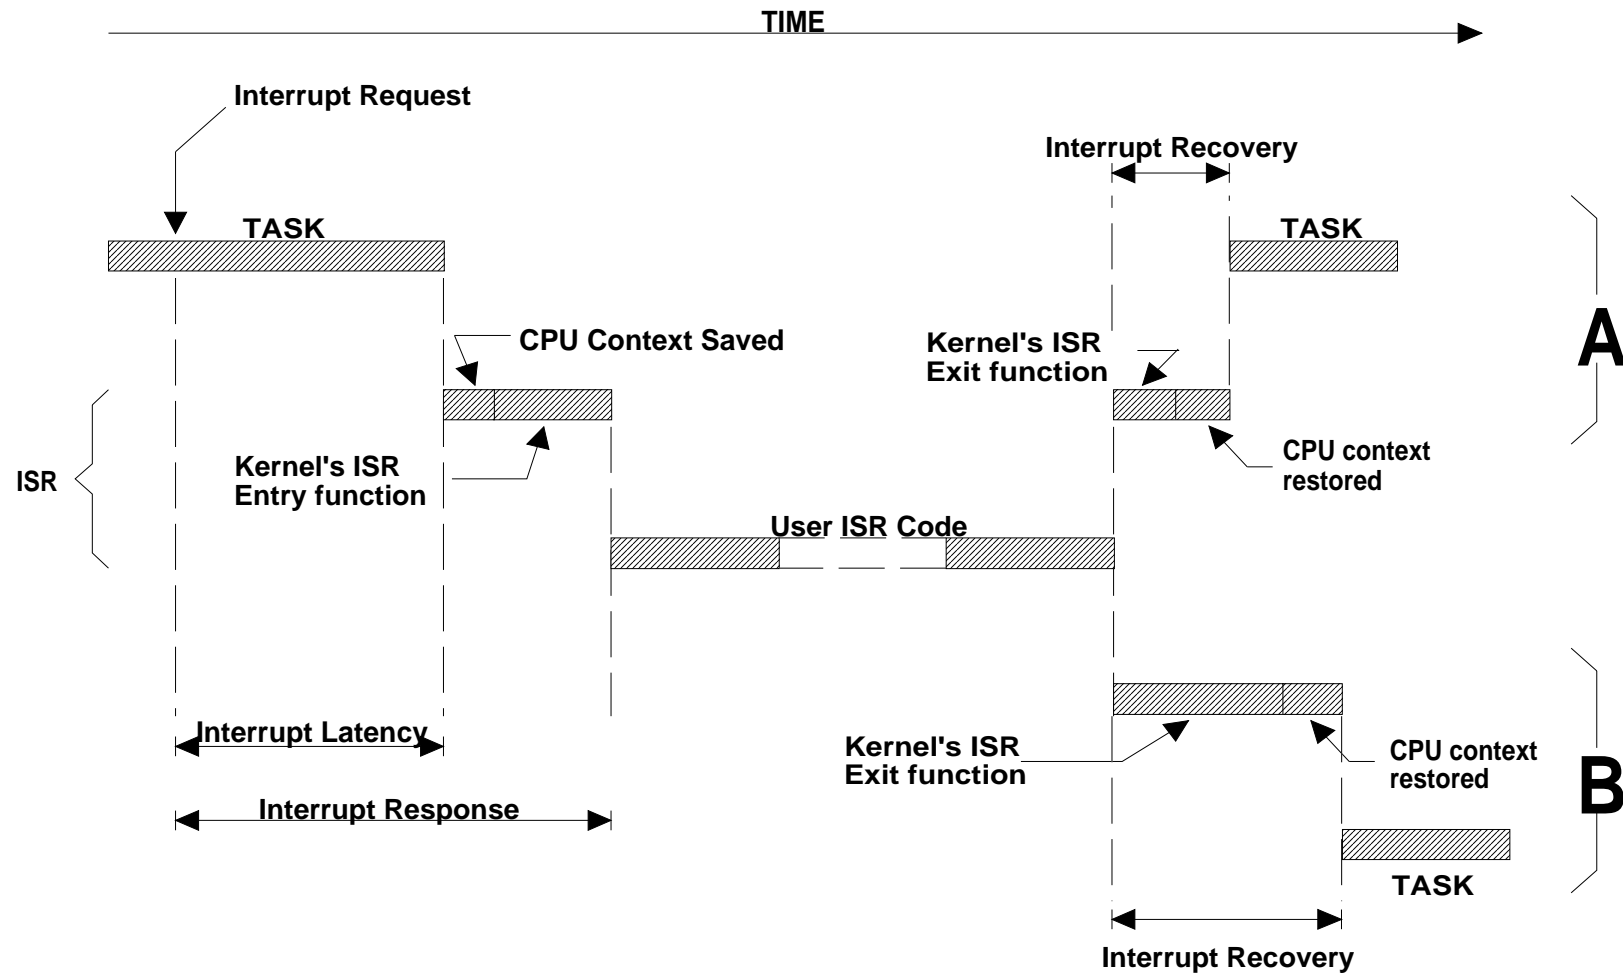
Interrupt

- ❑ An interrupt is a hardware mechanism used to inform the CPU that an asynchronous event has occurred.
- ❑ Interrupt Latency
- ❑ Interrupt Response
- ❑ Interrupt Recovery
- ❑ Interrupt Nesting
- ❑ Non-Maskable Interrupts (NMIs)

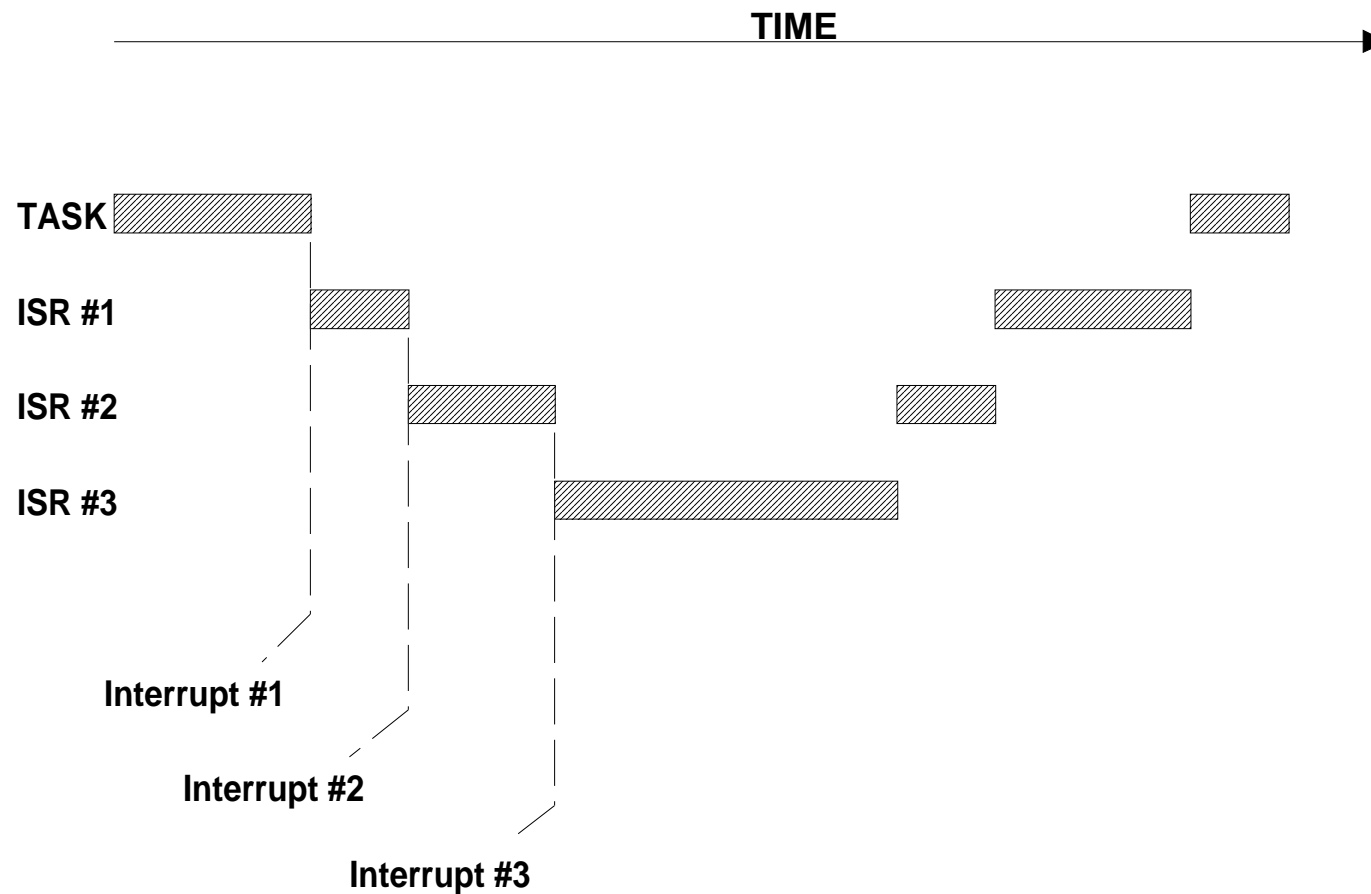
Non-preemptive kernel



Preemptive kernel



Interrupt Nesting



Non-Maskable Interrupts (NMIs)

- ❑ Service the most important time-critical ISR
 - Can not be disabled
- ❑ Interrupt latency = Time to execution the longest instruction + Time to start execution the NMI ISR
- ❑ Interrupt response = Interrupt latency + Time to save CPU context
- ❑ Interrupt recovery = Time to restore CPU context + time of executing return-from-interrupt

Clock Tick

- ❑ A periodical interrupt
 - Be viewed as heartbeat
 - Application specific and is generally between 10ms and 200ms
 - Faster timer causes higher overhead
 - Delay problem should be considered in real-time kernel.

Real-Time Kernels

- ❑ Real-time OS allows real-time application to be designed and expanded easily.
- ❑ The use of an RTOS simplifies the design process by splitting the application into separate tasks.
- ❑ Real-time kernel requires more ROM/RAM and 2 to 4 percent overhead.
 - Kernel request extra code(ROM)
 - Application code(RAM)
 - All kernel require extra RAM to maintain internal variables, data structure, queue, etc

Kernel Structure

Home Automation, Networking and Entertainment Lab.

CSIE @ National Cheng Kung University.



Kernel Structure

- ❑ Critical Sections
- ❑ Tasks
- ❑ Task Control Blocks
- ❑ Ready List
- ❑ Idle Task
- ❑ Statistic Task
- ❑ μ C/OS-II Initialization

Critical Sections

❑ OS_CRITICAL_METHOD == 1

- The first and simplest way to implement thus two macro is to invoke the processor instruction to disable interrupts for **OS_ENTER_CRITICAL()** and to enable interrupts instruction for **OS_EXIT_CRITICAL()**.

❑ OS_CRITICAL_METHOD == 2

- The second way to implement **OS_ENTER_CRITICAL()** is to save the interrupt disable status onto the stack and then disable interrupt. **OS_EXIT_CRITICAL()** is implement by restoring the interrupt status from stack.

Critical Sections(Cont.)

- ❑ OS_CRITICAL_METHOD == 3
 - Some compiler assume you with extensions that you to obtain the current value of Processor Status Word (PSW) and save it into a local variable within a C function.

Tasks

- ❑ Up to 64 tasks.
- ❑ Two tasks for system use (idle and statistic).
- ❑ Priorities 0, 1, 2, 3, OS_LOWEST_PRIO-3, OS_LOWEST_PRIO-2, OS_LOWEST_PRIO-1, OS_LOWEST_PRIO for future use.
- ❑ Each task must be assigned a unique priority level from 0 to OS_LOWEST_PRIO-2.
- ❑ The task priority is also the task identifier.

Multitasking

- ❑ Multitasking is started by calling OSStart().
- ❑ OSStart() runs the highest priority task that is **READY** to run.

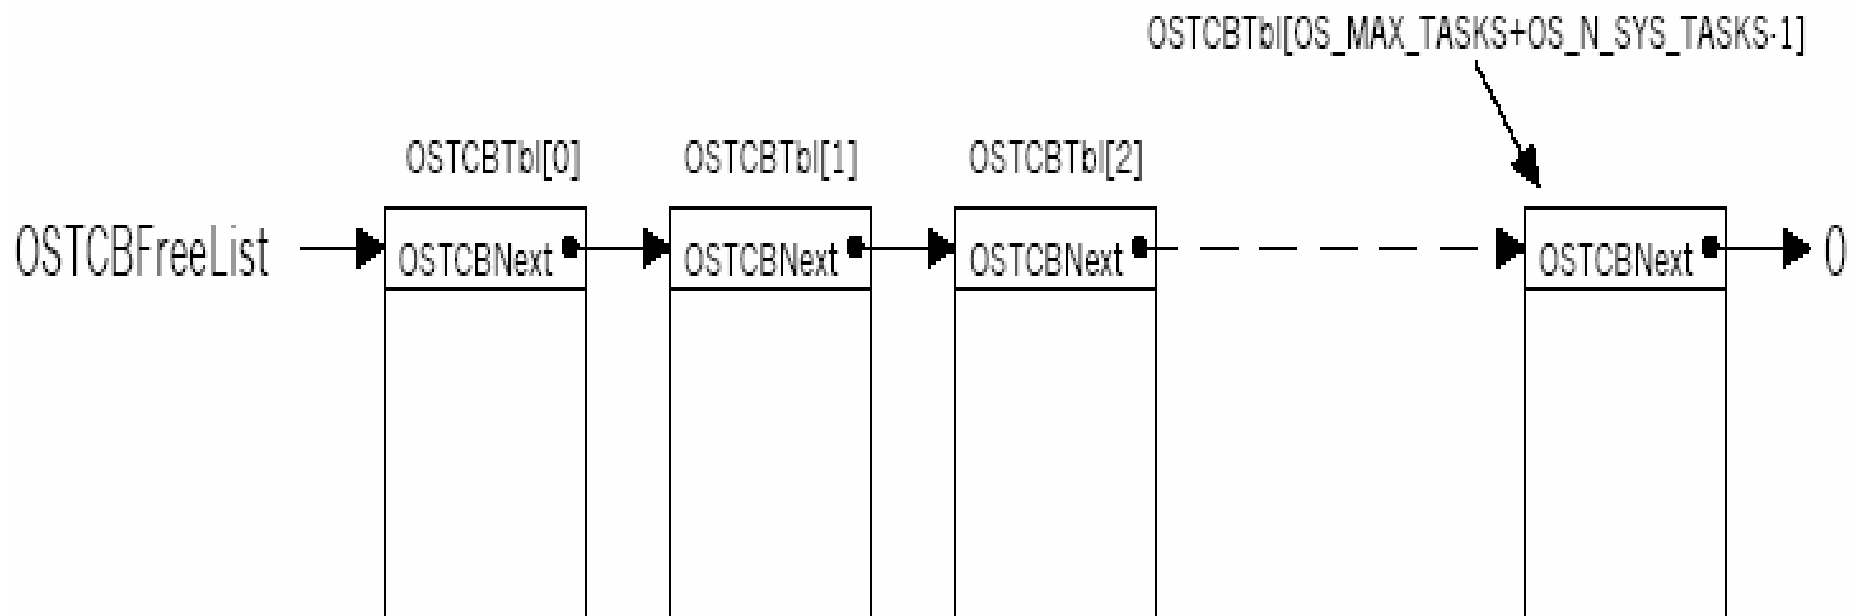
Task Control Blocks (Cont.)

- ❑ Used to accelerate the process of making a task ready to run, or to make a task wait for an event (to avoid computing these values at runtime).
- ❑ The values for these fields are computed when the task is created or when the task's priority is changed.

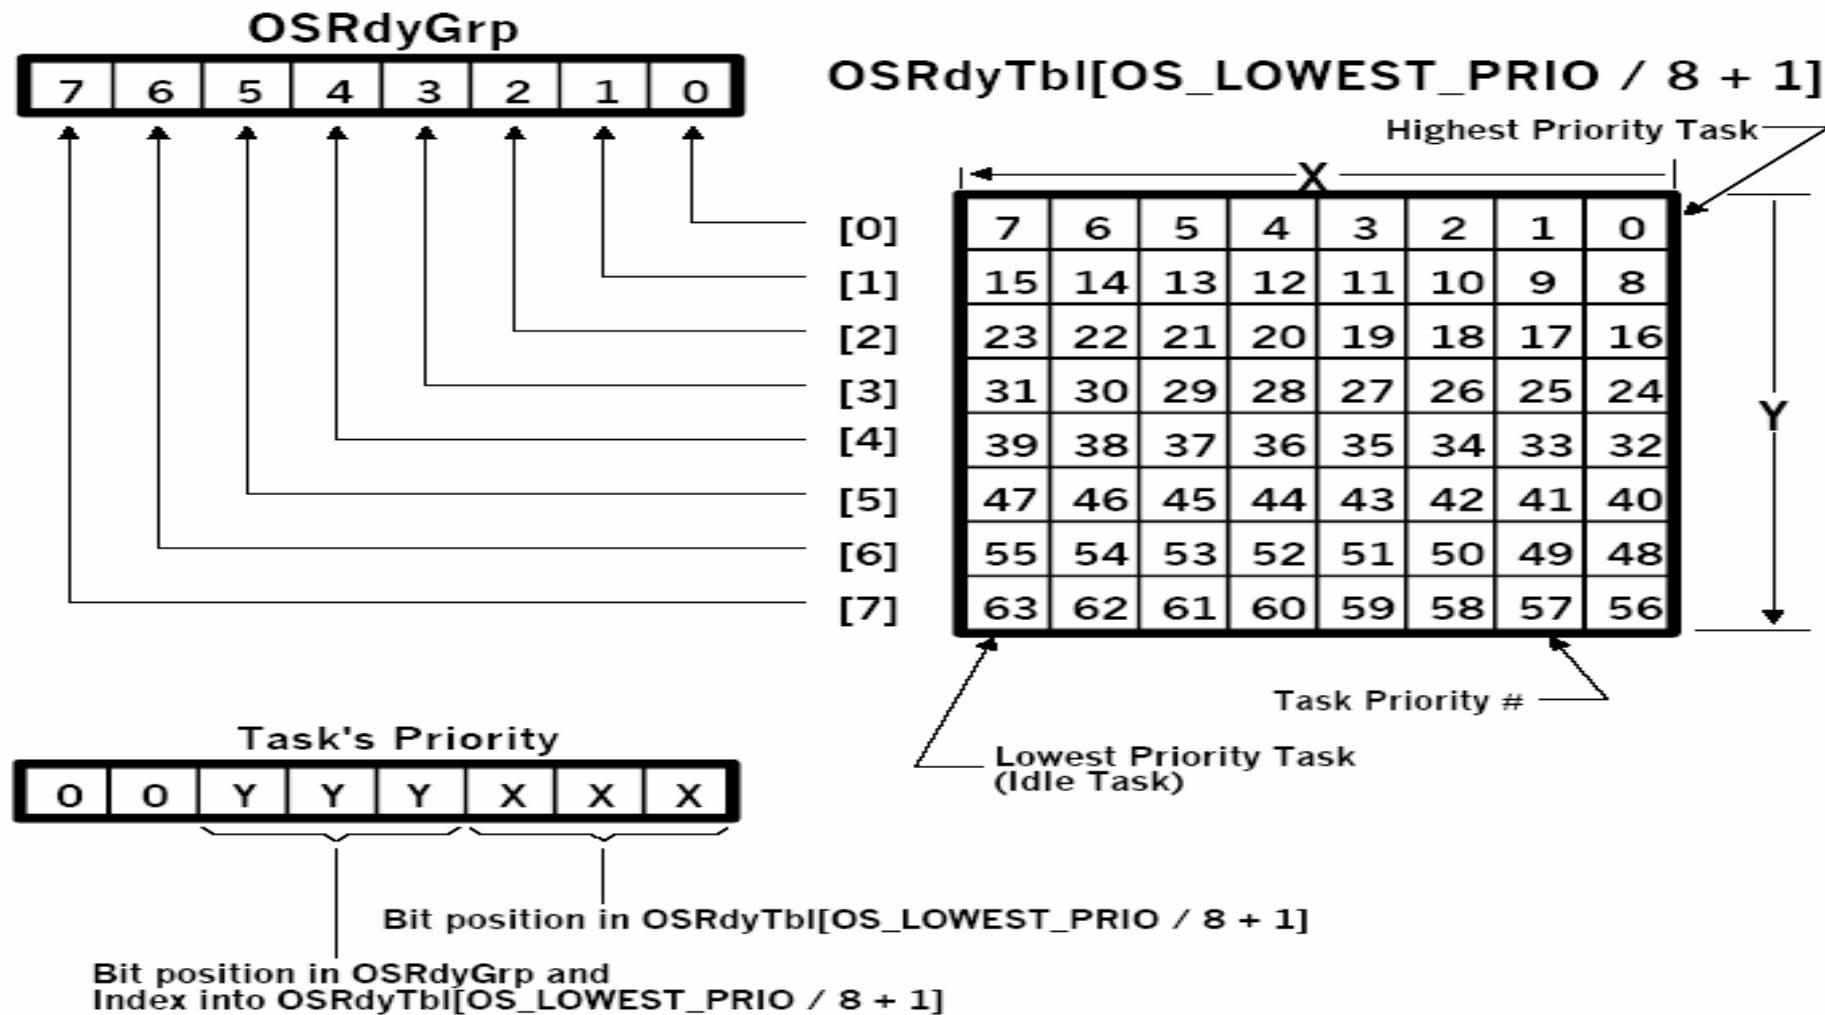
```
OSTCBY      = priority >> 3;  
OSTCBBitY   = OSMaTbl[priority >> 3];  
OSTCBX      = priority & 0x07;  
OSTCBBitX   = OSMaTbl[priority & 0x07];
```

OS_TCB Lists

- ❑ TCBs store in OSTCBTbl[].
- ❑ All TCBs are initialized and linked when uC/OS-II is initialized.



Ready List



Ready List(Cont.)

- ❑ Each task that is ready to run is placed in a ready list consisting of two variables, **OSRdyGrp** and **OSRdyTbl[]**.
- ❑ Task priorities are grouped (8 tasks per group) in OSRdyGrp.
- ❑ Each bit in OSRdyGrp is used to indicate whenever any task in a group is ready to run.
- ❑ When a task is ready to run it also sets its corresponding bit in the ready table, OSRdyTbl[].

OSRdyGrp and OSRdyTbl[]

- ❑ To determine which priority (and thus which task) will run next, the scheduler determines the lowest priority number that has its bit set in OSRdyTbl[].

Bit 0 in OSRdyGrp is 1 when any bit in OSRdyTbl [0] is 1.

Bit 1 in OSRdyGrp is 1 when any bit in OSRdyTbl [1] is 1.

Bit 2 in OSRdyGrp is 1 when any bit in OSRdyTbl [2] is 1.

Bit 3 in OSRdyGrp is 1 when any bit in OSRdyTbl [3] is 1.

Bit 4 in OSRdyGrp is 1 when any bit in OSRdyTbl [4] is 1.

Bit 5 in OSRdyGrp is 1 when any bit in OSRdyTbl [5] is 1.

Bit 6 in OSRdyGrp is 1 when any bit in OSRdyTbl [6] is 1.

Bit 7 in OSRdyGrp is 1 when any bit in OSRdyTbl [7] is 1.

Making a task ready to run

```
OSRdyGrp      |= OSMaTbl[prio >> 3];  
OSRdyTbl[prio >> 3] |= OSMaTbl[prio & 0x07];
```

Index	Bit Mask (Binary)
0	0000000 1
1	000000 1 0
2	00000 1 00
3	0000 1 000
4	000 1 0000
5	00 1 00000
6	0 1 000000
7	1 0000000

Task's Priority							
0	0	Y	Y	Y	X	X	X

Removing a task from the ready list

```
if ((OSRdyTbl[prio >> 3] &= ~OSMapTbl[prio & 0x07]) == 0)
    OSRdyGrp &= ~OSMapTbl[prio >> 3];
```

Finding the highest priority task

```
y      = OSUnMapTbl[OSRdyGrp];  
x      = OSUnMapTbl[OSRdyTbl[y]];  
prio = (y << 3) + x;
```

```
INT8U const OSUnMapTbl[] = {  
    0, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, 4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,  
    5, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, 4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,  
    6, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, 4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,  
    5, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, 4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,  
    7, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, 4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,  
    5, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, 4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,  
    6, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, 4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,  
    5, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, 4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0  
};
```

76543210	OSUnMapTbl[]
00000000	0
00000001	0
00000010	1
00000011	0
00000100	2
00000101	0
00000110	1
00000111	0
00001000	3
00001001	0
00001010	1
00001011	0
00001100	2
00001101	0
00001110	1
00001111	0
...	

Idle Task

- ❑ μ C/OS-II always creates a task (a.k.a. the *Idle Task*) which is executed when none of the other tasks is ready to run.
- ❑ The idle task (OSTaskIdle()) is always set to the lowest priority, i.e. OS_LOWEST_PRIO. OSTaskIdle() does nothing but increment a 32-bit counter called OSIdleCtr.
- ❑ Interrupts are disabled then enabled around the increment because on 8-bit and most 16-bit processors, a 32-bit increment requires **multiple instructions** which must be protected from being accessed by higher priority tasks or an ISR.

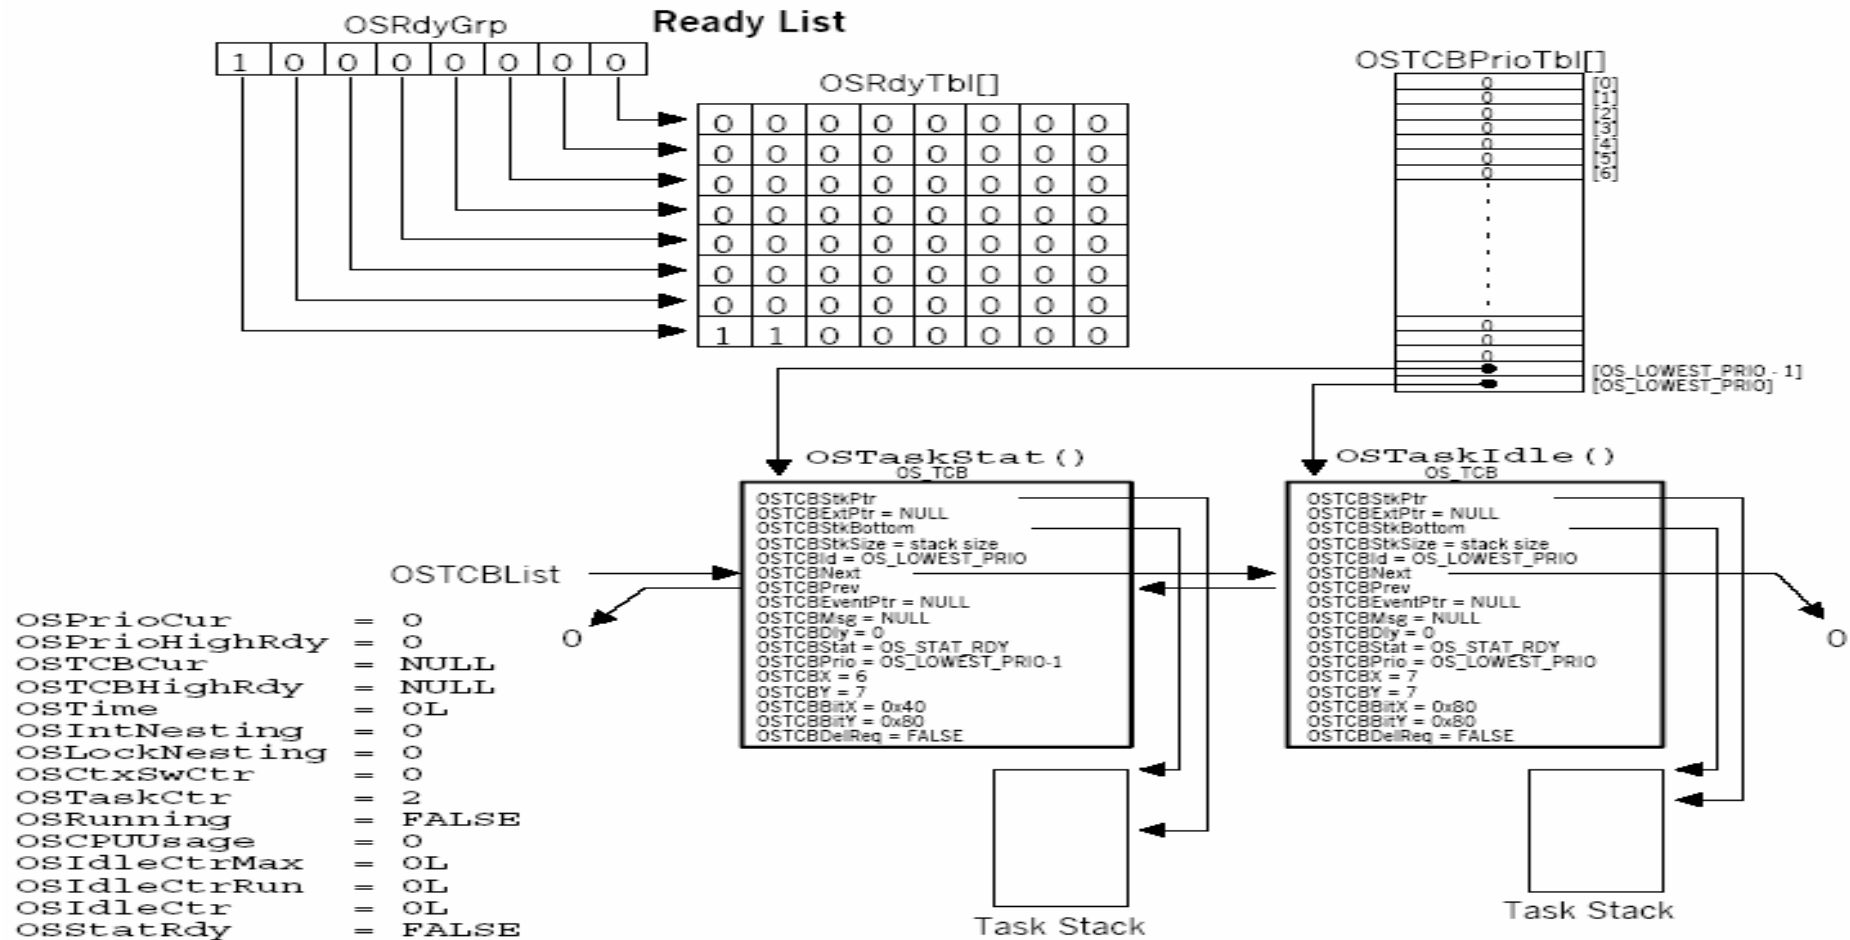
Statistics Task

- ❑ μ C/OS-II contains a task that provides run-time statistics.
- ❑ OSTaskStat() executes every second and computes the percentage of CPU usage.
- ❑ This value is placed in the variable OSCPUUsage which is a signed 8-bit integer.
- ❑ You MUST call OSStatInit() from the first and only task created in your application during initialization.

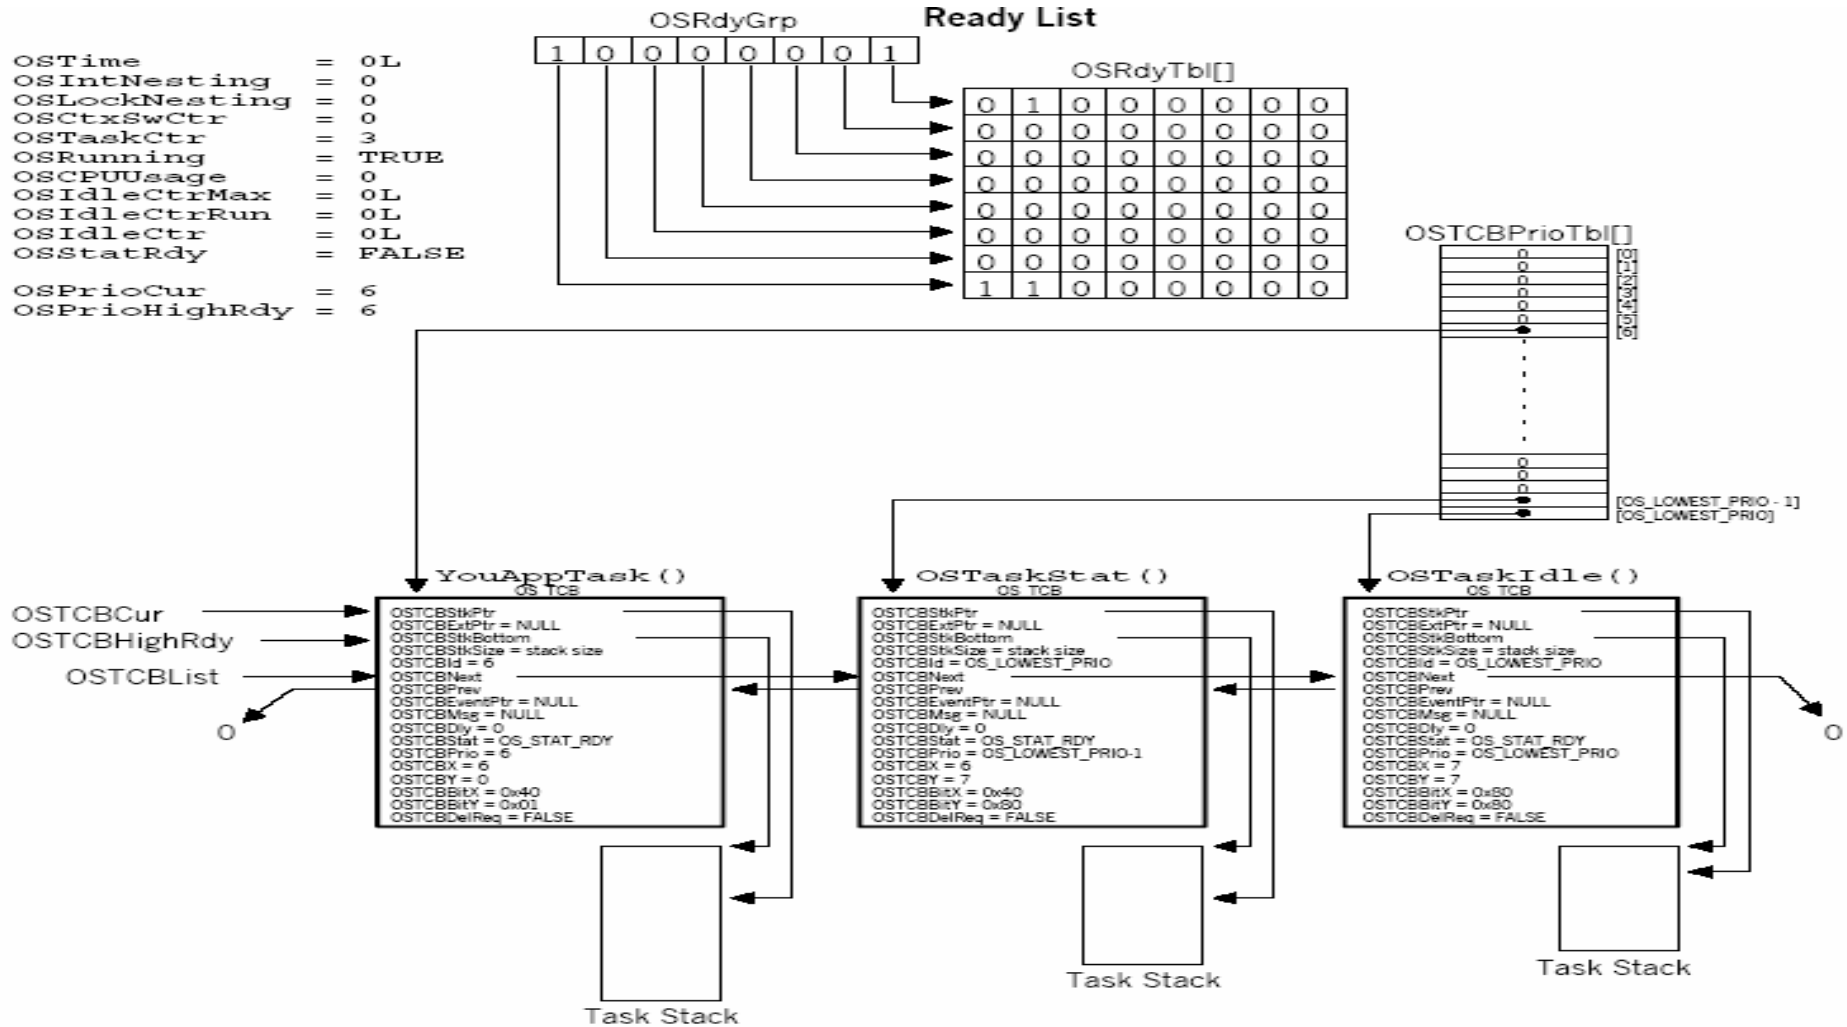
uC/OS-II Initialization

- ❑ A requirement of μ C/OS-II is that you call OSInit() before you call any of its other services.
- ❑ OSInit() initializes all of μ C/OS-II's variables and data structures.
- ❑ Create the idle task OSTaskIdle()
- ❑ Create the statistic task OSTaskStat()
- ❑ The Task Control Blocks (OS_TCBs) of these two tasks are chained together in a doubly-linked list.

Data structures after calling OSInit()



Variables and Data Structures after calling OSStart()



Porting MicroC/OS-II

Home Automation, Networking and Entertainment Lab.

CSIE @ National Cheng Kung University.



Porting MicroC/OS-II

- ❑ Porting MicroC/OS-II to the x86 protected mode
- ❑ Loading the MicroC/OS-II
- ❑ Initializing the hardware
- ❑ Converting to a 32-bits, flat memory model
- ❑ Building the application

Porting MicroC/OS-II to the x86 protected mode

- ❑ Porting MicroC/OS-II to a protected mode, 32-bit, flat memory model of the 80386, 80486, Pentium and Pentium II CPUs.
- ❑ MicroC/OS-II is a portable, ROMable, preemptive, real-time, multitasking kernel for microprocessors.
- ❑ Originally written in C and x86 (real mode) assembly languages, it has been ported to a variety of microprocessors.

Loading the MicroC/OS-II

- ❑ The MicroC/OS-II can be loaded in various ways:
 - a) It can be loaded from DOS. But switching into protected mode and fully bypassing DOS and the BIOS may create some inconsistencies inside DOS, preventing a normal return to it. Thus, the PC would have to be rebooted in order to get back to DOS. Also, the application would have to be built as a DOS application, requiring a 16-bit compiler.

Loading the MicroC/OS-II(Cont.)

- b) It can be loaded from a floppy disk upon boot-up. By providing a bootstrap loader, the application would be the first thing loaded and would be in full control. The PC also has to be rebooted back to DOS (or Windows) if required.
- c) It can be burned into PROM, into a stand-alone system.

Bootstrap Loader

- ❑ The bootstrap loader (BootSctr.img) can easily be installed on the very first sector of a floppy disk (e.g. the boot sector) by using the DEBUG program, distributed with DOS and Windows.
- ❑ The DEBUG's write command has the following format:
 - *-w offset disk track sector_count*
- ❑ By executing the following commands:
 - C:\MyTask>debug bootstrap.img
 - -w 0100 0 0 1
 - -q

Bootstrap Loader(Cont.)

- ❑ When loaded, this bootstrap loader:
 - a) Loads the first 64k of the floppy disk (the first file in fact) at the physical address 1000h.
 - b) Disables the interrupts.
 - c) Jumps at 1000h to start executing the application.

Initializing the Hardware(Cont.)

- ❑ Once in *main()*, a call is done to *OsCpuInit()*, in *os_cpu_c.c*, in order to perform the following:
 - Enable the address line 20, normally disabled for some real mode considerations. The line is enabled by sending a few commands to the Intel 8042 keyboard controller. See *InitA20()* for details (*os_cpu_c.c*).
 - Relocate the IRQ interrupts, since the overlap the CPU interrupts and exceptions (for instance, it is not possible to know if the interrupt 0 has been triggered by the clock or a division by zero). The IRQ are relocated in the range 20h-2Fh by sending a few commands to the Intel 8259 interrupt controllers. See *InitPIC()* for details (*os_cpu_c.c*).

Initializing the Hardware(Cont.)

- The interrupt table is initialized by using *SetIntVector()* and *SetIDTGate()*. The 64 entries are set to point to a default interrupt handler (*DefIntHandler()*, in *os_cpu_a.asm*), which simply performs an interrupt return.
- The clock handler (*OsTickISR()*, in *os_cpu_a.asm*) is installed as the interrupt 20h handler.
- The MicroC/OS context switch handler (*OSCtxSw()*, in *os_cpu_a.asm*), is installed as the interrupt *uCOS* handler (*uCOS* is defined as 0x30 in *os_cpu.h*).